

Output: If the code executes without errors, determine the output.
Concept: Briefly state the key programming concepts the question tests if provided in the options.
Review the Code below and determine from amongst the provided options

```
1 | #include <stdio.h>
2 | #include <string.h>
3 | int main() {
4 |     char *str1 = "hello";
5 |     char *str2 = "world";
6 |     str1 = str2;
7 |     printf("%s\n", str1);
8 |     return 0;
9 | }
```

Answer:

- ☐ The code outputs "hello"
- ☐ The code outputs "world"
- ☐ Compiler error: Cannot modify a string literal
- ☐ Runtime Error

[< Previous](#)

[Clear Answer](#)




```
2 | void reverse_array(int *arr, int start, int end) {  
3 | while (start <= end) {  
4 | int temp = arr[start];  
5 | arr[start] = arr[end];  
6 | arr[end] = temp;  
7 | start++;  
8 | end--;  
9 | }  
10 | }  
11 | int main() {  
12 | int numbers[] = {1, 4, 2, 8, 5};  
13 | reverse_array(numbers, 0, 4);  
14 | for (int i = 0; i < 5; i++) {  
15 | printf("%d ", numbers[i]);  
16 | }  
17 | printf("\n");  
18 | return 0;  
19 | }
```

Answer:

- ☐ Change the comparison in the while loop to <.
- ☐ Initialize start to 1 in main.
- ☐ The code functions correctly as intended.
- ☐ Use pointers instead of indexing in reverse_array.

< Previous

Clear Answer



```
1 | #include <string.h>
2 |
3 | char* replace_substring(char* str, const char* old_str, const
4 | char* new_str) {
5 |     char *result;
6 |     char *pos;
7 |     pos = strstr(str, old_str);
8 |     if (pos == NULL) {
9 |         return str;
10 |    }
11 |    int new_len = strlen(new_str);
12 |    int old_len = strlen(old_str);
13 |    result = (char*)malloc(strlen(str) + new_len - old_len + 1);
14 |    strncpy(result, str, pos - str);
15 |    strcpy(result + (pos - str), new_str);
16 |    strcpy(result + (pos - str) + new_len, pos + old_len);
17 |    return result;
18 | }
```

Answer:

- ☐ Change the allocation size in malloc to account for potential differences in the length of the old and new substrings.
- ☐ Add a loop to handle multiple occurrences of old_str.
- ☐ The code is perfectly correct.
- ☐ Use memmove instead of strcpy for overlapping memory regions.

[< Previous](#)

[Clear Answer](#)





```
1 | #include <stdio.h>
2 | #include <math.h>
3 | int is_prime(int num)
4 | {
5 |     if (num <= 1) {
6 |         return 0;
7 |     }
8 |     for (int i = 2; i < num; i++) {
9 |         if (num % i == 0) {
10 |             return 0;
11 |         }
12 |     }
13 |     return 1;
14 | }
```

Answer:

- ☐ The code has no errors.
- ☐ Change the initial value of i to 3.
- ☐ Add a special check for the number 2.
- ☐ Change the for loop condition to $i \leq \sqrt{\text{num}}$.

[< Previous](#)

[Clear Answer](#)



```
1 | #include <stdio.h>
2 | int find_max(int arr[], int size)
3 | {
4 |     int max = 0;
5 |     for (int i = 0; i < size; i++) {
6 |         if (arr[i] > max) {
7 |             max = arr[i];
8 |         }
9 |     }
10 |    return max;
11 | }
12 | int main() {
13 |     int numbers[] = {3, 7, 1, 9, 5};
14 |     int max_value = find_max(numbers, 5);
15 |     printf("Max value: %d\n", max_value);
16 |     return 0;
17 | }
```

Answer:

- ☐ Initialize max to the first element of the array: `int max = arr[0];`
- ☐ Change the array access in the loop to `arr[i + 1]`.
- ☐ Return a pointer to the max element within the array.
- ☐ The code is perfectly correct.

[< Previous](#)

[Clear Answer](#)



What will be the output of the following Program:

```
1 | #include <stdio.h>
2 | void fun(int*, int*);
3 | int main()
4 | {
5 |     int i=5, j=2;
6 |     fun(&i, &j);
7 |     printf("%d, %d", i, j);
8 |     return 0;
9 | }
10 | void fun(int *i, int *j)
11 | {
12 |     *i = *i++;
13 |     *j = *j++;
14 | }
```

Answer:

- ☐ 10, 4
- ☐ 5, 2
- ☐ 25, 4
- ☐ 2, 5

[< Previous](#)

[Clear Answer](#)



Concept: Briefly state the key programming concepts the question tests if provided in the options
Choose the Correct Option:

```
1 | #include <string>
2 | #include <iostream>
3 | using namespace std;
4 | int main() {
5 |     string str = "hello";
6 |     str[5] = '!';
7 |     cout << str << endl;
8 |     return 0;
9 | }
```

Answer:

- ☐ The code outputs "hello!"
- ☐ The code outputs "hell!"
- ☐ Compiler error: Invalid string initialization
- ☐ Runtime Error: Accessing an out-of-bounds array index

[< Previous](#)[Clear Answer](#)

Concept: Briefly state the key programming concepts the question tests, provided in the options.
For the code snippet below, identify the most appropriate correction or improvement from the given options.


```
1 | #include <stdio.h>
2 | #include <stdlib.h>
3 | int main() {
4 |     int *ptr = malloc(sizeof(int));
5 |     if (ptr == NULL) {
6 |         printf("Memory allocation failed.\n");
7 |     }
8 |     *ptr = 20;
9 |     // ... more code using ptr
10 | return 0;
11 | }
```

Answer:

- ☐ Use realloc instead of malloc.
- ☐ Cast the return value of malloc: `int * ptr = (int *)malloc(sizeof(int));`
- ☐ The code is correct as is.
- ☐ Add a call to `free(ptr)` at the end to release memory.

[< Previous](#)

[Clear Answer](#)




```
1 | #include <iostream>
2 | using namespace std;
3 | void increment(int val) {
4 |     val++;
5 | }
6 | int main() {
7 |     int x = 10;
8 |     increment(x);
9 |     cout << x << endl;
10 | return 0;
11 | }
```

Answer:

- ☐ Compiler error: Missing return statement in 'increment'
- ☐ The code outputs 11
- ☐ The code compiles, but has no output
- ☐ The code outputs 10

[< Previous](#)

[Clear Answer](#)



Question: 50

To address the "convoy effect" sometimes seen in SJF or priority-based scheduling, a potential solution is:

Answer:

- ☐ Aging – gradually increasing priority of waiting jobs.
- ☐ Strict real-time deadlines with admission control.
- ☐ Periodically reducing the priority of long processes.
- ☐ Using a lottery system for selecting the next job.

[< Previous](#)

[Clear Answer](#)



Question: 49

Rotational Positional Sensing (RPS) is an optimization primarily targeting:

Answer:

- ☐ Minimizing power consumption of spinning disks.
- ☐ Improving effective throughput over zoned-bit recording disks.
- ☐ Reducing average rotational delay to read a sector.
- ☐ Fault tolerance in distributed storage systems.

[< Previous](#)

[Clear Answer](#)



Review the Code below and determine from amongst the provided options

```
1 | #include <stdio.h>
2 | struct Point {
3 |     int x;
4 |     int y;
5 | };
6 | int main() {
7 |     struct Point p1 = {10, 20};
8 |     struct Point *ptr = &p1;
9 |     printf("%d\n", *ptr->x);
10 | return 0;
11 | }
```

Answer:

- ☐ Compiler error: Invalid struct initialization
- ☐ Runtime error
- ☐ The code outputs 10
- ☐ The code outputs 20

[< Previous](#)

[Clear Answer](#)



Question: 41

Consider a cache with 64-byte cache lines, 2-way set-associativity and a 32-bit physical address space. What's the total number of cache entries? (Assume 4 KB of Cache Size)

Answer:

- ☐ 26
- ☐ 6
- ☐ 21
- ☐ 19

< Previous

Clear Answer



Question: 53

Which of the following best describes the concept of eventual consistency in NoSQL databases?

Answer:

- ☐ Data consistency is sacrificed entirely in favor of high availability.
- ☐ Data updates only become visible after a commit log is flushed to disk.
- ☐ Data consistency will be achieved at some point in the future, but immediate reads may return stale data.
- ☐ Data consistency is enforced immediately across all replicas.

[< Previous](#)

[Clear Answer](#)



Question: 55

Which data model is most commonly associated with MongoDB?

Answer:

- ☐ Document
- ☐ Graph
- ☐ Key-Value
- ☐ Column-oriented

[< Previous](#)

[Clear Answer](#)



Question: 54

Consider a scenario where you need to perform sentiment analysis on millions of social media posts. Which Big Data technology stack would be best for this task?

Answer:

- ☐ Apache Storm, Spark Streaming, Machine Learning Library
- ☐ Hadoop, Hive, Pig
- ☐ Spark, Kafka, Elasticsearch
- ☐ HBase, Cassandra, Zookeeper

Clear Answer

< Previous



Question: 44

The readers-writers problem prioritizes read access when possible. Implementing a "reader priority" solution risks violating which condition?

Answer:

- ☐ Bounded Waiting for writers
- ☐ Progress for readers
- ☐ No reader kept waiting unless a writer has permission.
- ☐ Mutual Exclusion for writers

[< Previous](#)[Clear Answer](#)

Question: 42

Why is interleaved memory often preferred in high-performance systems?

Answer:

- ☐ It permits simultaneous access to consecutive memory locations.
- ☐ It prevents memory bank conflicts, reducing bus contention.
- ☐ It ensures data is distributed for memory-level parallelism.
- ☐ It eliminates the need for virtual memory support.

[< Previous](#)

[Clear Answer](#)



Question: 45

Why is interrupt-driven I/O generally favoured over programmed I/O (polling)?

Answer:

- ☐ It guarantees immediate servicing of high-priority devices.
- ☐ It eliminates the need for hardware synchronization of I/O devices.
- ☐ It requires less complex programming support from the OS.
- ☐ It allows the CPU to overlap other work while I/O occurs.

[< Previous](#)

[Clear Answer](#)



Question: 46

Modern CPUs often provide atomic instructions like 'test-and-set' or 'compare-and-swap'. A primary use case for these within an OS is

Answer:

- ☐ Context switching with reduced register saving overhead.
- ☐ Handling device interrupts safely.
- ☐ Implementing message passing between processes.
- ☐ Building low-level spinlocks for fine-grained synchronization.

[< Previous](#)[Clear Answer](#)

Question: 48


Which factor has the LEAST impact when choosing a disk scheduling algorithm for an interactive general-purpose workstation?

Answer:

- ☐ Maximizing sequential read throughput for large files
- ☐ Responsiveness to short file read/write requests
- ☐ Minimizing average seek time
- ☐ Fairness of wait times across different processes

< Previous

Clear Answer



Question: 49

A system uses a multilevel feedback queue, but processes seem stuck at lower priority levels. A likely cause is

Answer:

- ☐ I/O operations aren't releasing the CPU as expected.
- ☐ The time quantum at the upper priority level is too short.
- ☐ The time quantum at the upper priority level is too long.
- ☐ There are not enough CPU resources for the workload.

[< Previous](#)[Clear Answer](#)

Question: 22

You need an algorithm to determine if there exists a path between two vertices in a directed graph. Which is the most suitable choice?

Answer:

- ☐ Prim's algorithm
- ☐ Any graph traversal algorithm (DFS or BFS)
- ☐ Dijkstra's Algorithm
- ☐ Kruskal's algorithm

< Previous

Clear Answer



Question: 36

You need a self-balancing binary search tree, which guarantees worst-case $O(\log n)$ operations at the expense of potentially complex rotations.

Answer:

- ☐ Red-black tree
- ☐ AVL Tree
- ☐ Splay tree
- ☐ B-tree

[< Previous](#)[Clear Answer](#)

Question: 47

In a paged system, the Page Table Base Register (PTBR) is crucial. When a process switch occurs, which action MUST be taken?

Answer:

- ☐ All of the mentioned
- ☐ The CPU must enter privileged mode to access the PTBR.
- ☐ All entries in the TLB must be invalidated.
- ☐ The PTBR must be updated to point to the new process's page table.

[< Previous](#)

[Clear Answer](#)



Question: 32

For pattern matching where the pattern may contain wildcards or special characters, which algorithm outperforms simple brute-force?

Answer:

- ☐ Knuth-Morris-Pratt (KMP) algorithm
- ☐ Naive string search
- ☐ Aho-Corasick algorithm
- ☐ Rabin-Karp algorithm

[< Previous](#)[Clear Answer](#)

Question: 25

What data structure underlies the implementation of most version control systems like Git for tracking changes to a file over time?

Answer:

- ☐ Directed acyclic graph (DAG)
- ☐ Bloom filter
- ☐ Merkle tree
- ☐ Linked list

[< Previous](#)[Clear Answer](#)

Question: 247

Which algorithm would you use to generate all possible permutations of a given array?

Answer:

- ☐ Depth-first search (with backtracking)
- ☐ Merge sort
- ☐ Heap's algorithm
- ☐ Topological sort

< Previous

Clear Answer



DBMS

a) integration and analysis - - -

b) Time varied attributes

c) Data consistency is sacrificed

54) - Apache Storm, Spark Streaming

55) Docucl

56) Truesform

57) Batch process of logs, unstructured

58) consistency, Availability, Partition Tolerance

59)

60)

~~Fixed schema over and over need~~
~~horizontal scaling~~
Fixed schema

Question: 33

You need a data structure for a real-time leader board with these operations:

insertScore(player, score), deleteScore(player), and getTopK(K). Which is the MOST efficient combination?

Answer:

- ☐ Skip list for insertion/deletion, min-heap for getTopK.
- ☐ Self-balancing BST (e.g., AVL tree) for insertion, hash table for deletion.
- ☐ Max-heap for insertion/lookup, array for deletion.
- ☐ Sorted array for insertion, hash table for lookup.

< Previous

Clear Answer



Answer:

- ☐ Quantitative measures and metrics
- ☐ Time-variant attributes that describe entities
- ☐ Primary and foreign keys
- ☐ Textual descriptions of entities

[< Previous](#)

[Clear Answer](#)



Question: 40

What kind of tree allows constant-time access to the minimum AND maximum element within it?

Answer:

- ☐ Binary heap
- ☐ Binary search tree (unmodified)
- ☐ AVL tree
- ☐ Cartesian Tree

< Previous

Clear Answer



Question: 38

There exists an algorithm with time complexity $O(n)$ to solve a problem. Processing an input of size $n=5$ takes 0.2 seconds. Estimate the MAXIMUM time you could reasonably process in under 1 minute.

Answer:

- ☐ $n=8$
- ☐ $n=9$
- ☐ $n=7$
- ☐ $n=6$

Clear Answer

< Previous



simulated annealing
dynamic programming
Adleman matrix

Trie

algorithm B (mergesort)

26) 0 (ms)

27) Encapsulation

28) Count min sketch

29) Rabin-Karp algorithm

30) polynomials

31) g midiset

32) Aho - Corasick algorithm

33)

34) Depth first search

35)

36)

Quicksort

37)

AVL tree

38)

DFS on binary tree

39)

$m = 6$

40)

CSR format

Binary heap

Question: 3

An algorithm demonstrates logarithmic growth. If processing 10 items takes 3 milliseconds, approximately how long would 1000 items take?

Answer:

- ☐ 300 milliseconds
- ☐ 100 milliseconds
- ☐ 9 milliseconds
- ☐ 30 milliseconds

< Previous

Clear Answer



Question: 20

Which principle is BEST demonstrated by providing multiple implementations of a 'sort()' method, tailored to different data types within

Answer:

- ☐ Polymorphism
- ☐ Abstraction
- ☐ Inheritance
- ☐ Encapsulation

[< Previous](#)[Clear Answer](#)

Question: 25

Heavy use of the protected access modifier within a class hierarchy might signal a potential violation of which OOP principle?

Answer:

- ☐ Encapsulation
- ☐ Polymorphism
- ☐ Inheritance
- ☐ Liskov Substitution Principle

[< Previous](#)[Clear Answer](#)

Question: 29

To search for an extremely rare pattern in a massive text corpus, with failure being common but early rejection of non-matches desirable, which algorithm is most suitable?

Answer:

- ☐ Knuth-Morris-Pratt Algorithm
- ☐ Rabin-Karp Algorithm
- ☐ Boyer-Moore Algorithm
- ☐ Naive string search

< Previous

Clear Answer



Question: 26

To track the real-time frequency of words in a streaming text feed, which data structure keeps memory usage reasonable while allowing

Answer:

- ☐ Suffix array
- ☐ Min-heap
- ☐ Self-balancing binary search tree
- ☐ Count-min sketch

[< Previous](#)[Clear Answer](#)

Question 39

Implement a sparse matrix efficiently (many zero elements). Which is most memory-conscious but preserves reasonable element access?

Answer:

- ☐ Compressed sparse row (CSR) format
- ☐ Hash table mapping coordinates to values
- ☐ List of key-value pairs (coordinates and value)
- ☐ Store in a standard 2D array

[Clear Answer](#)[← Previous](#)

Question: 26

You have a recursive algorithm with the recurrence relation $T(n) = 2T(n/4) + n$. Using the Master Theorem, approximate its time complexity.

Answer:

- ☐ $O(n \log n)$
- ☐ $O(n \log^2 n)$
- ☐ $O(n^2)$
- ☐ $O(n \sqrt{n})$

Clear Answer

< Previous



Question: 25

Two sorting algorithms are considered:

Algorithm A: Best case $O(n)$, Average/Worst case $O(n^2)$

Algorithm B: All cases $O(n \log n)$

When sorting extremely large almost pre-sorted datasets, which would likely be faster?

Answer:

- ☐ Impossible to know without testing on the specific data
- ☐ Algorithm B (e.g., merge sort's consistency wins)
- ☐ They would perform roughly the same
- ☐ Algorithm A (e.g., insertion sort would excel)

< Previous

Clear Answer



Time Remaining: 02:05:20

CAMS

Question: 24

Need to efficiently store and retrieve IP addresses, supporting prefix-based lookups (e.g., find all addresses starting with "192.168.3"). Which data structure is optimal?

Answer:

- ☐ Binary search tree
- ☐ Hash table
- ☐ Bloom filter
- ☐ Trie (prefix tree)

Clear Answer

< Previous



Question: 20

9. Approximate the solution to the notoriously difficult Traveling Salesman Problem, which heuristic approach often yields decent results?

Answer:

- ☐ Simulated annealing
- ☐ 2-opt or 3-opt local search
- ☐ Branch and bound
- ☐ Backtracking

< Previous

Clear Answer



Question: 23

Which graph representation is generally best for a dense graph (where most potential edges exist)?

Answer:

- ☐ Incidence matrix
- ☐ Adjacency matrix
- ☐ Edge list
- ☐ Adjacency list

[< Previous](#)

[Clear Answer](#)



Question: 22

A computationally expensive optimization problem seems intractable to compromise, you consider an approach that rapidly produces good but not necessarily optimal solutions. This describes

Answer:

- ☐ Divide and Conquer
- ☐ Dynamic Programming
- ☐ Approximation Algorithm
- ☐ Greedy Algorithm

[Clear Answer](#)[< Previous](#)